
CAD-bench: Benchmarking Language Models on Functional CAD Generation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Can a language-model agent make a CAD artifact that actually satisfies a mechan-
2 ical task? Existing evaluations often stop after the code runs, the model renders,
3 or its surface roughly matches a reference. Those checks can miss a hole in the
4 wrong place, incompatible threads, colliding gears, or a transmission that does not
5 move. We introduce CAD-bench, an executable benchmark with 17 tasks ranging
6 from basic solids to threaded pairs and simulated gear trains. A system produces
7 CAD code or a STEP file; the grader then checks the measurements named in
8 the prompt, mating geometry, and, for gearboxes, motion in Blender. The best
9 of 19 complete standalone-model runs scores 59.9%, and the best of 32 complete
10 agent-harness runs scores 67.7%. The main gap is functional CAD: standalone
11 runs average 3.7% and agent runs 14.1% on the four functional tasks, even though
12 many of those submissions build. Across all 867 scored task attempts, 108 artifacts
13 build but receive zero or floor-level credit. Scorer tests also expose limitations:
14 independently expressed valid alternatives average 89.6%, including one false
15 rejection caused by a reference-shape gate, while several small faults retain high
16 scores. CAD-bench is therefore a diagnostic benchmark, not a claim of complete
17 CAD coverage: it shows where build success overstates mechanical correctness
18 and makes both grader successes and blind spots inspectable.

19 1 Introduction

20 CAD is more than a picture of an object. It records exact sizes, feature locations, mating surfaces,
21 and assemblies that later tools use for inspection, simulation, and manufacturing. Language-model
22 agents can now write CAD code and revise it after tool feedback, but evaluating the resulting artifact
23 remains difficult.

24 A generated artifact can compile and look plausible while still failing its task. A hole may be shifted,
25 two threads may not mate, or a pair of gears may intersect instead of transmitting motion. CAD-bench
26 is designed around this gap. It asks whether the submitted artifact satisfies the measurements and
27 mechanical behavior stated in the task, not whether it merely resembles a reference.

28 Prior CAD-generation benchmarks have established important foundations. SketchGraphs repre-
29 sents CAD sketches as geometric constraint graphs [Seff et al., 2020]. CAD-as-Language and
30 DeepCAD model CAD sketches or construction histories as serialized design programs [Ganin
31 et al., 2021, Wu et al., 2021]. Fusion 360 Gallery provides human-authored design sequences and a
32 programmatic reconstruction environment [Willis et al., 2021]. More recent systems move toward
33 natural-language CAD generation, executable CAD code, or geometric validation. Text2CAD studies
34 text-to-parametric-CAD generation from natural-language prompts [Khan et al., 2024]. CADPrompt
35 pairs natural-language CAD prompts with expert CAD code and uses a visual verification loop
36 [Alrashedy et al., 2024]. CAD-Coder and Text-to-CadQuery target text-to-CAD-code generation




Figure 1: CAD-bench evaluation pipeline. A harness receives a task specification and produces CAD code or a STEP artifact. The runtime executes the submission, exports geometry, and applies task-specific checks for dimensions, pose, thread observables, mating behavior, and rigid-body function before aggregation.

with executable feedback or geometric rewards [Guan et al., 2025, Xie and Ju, 2025]. CadEval, CAD-Smith, and CAD Arena provide related evaluation settings based on rendering, geometry, validity, or programmatic measurement [Epoch AI, 2026, Barkley et al., 2026, CAD Arena, 2026]. These benchmarks and systems show that executable geometric feedback is valuable. CAD-bench, however, targets a complementary and more significant question: whether a submitted CAD artifact satisfies mechanically meaningful task requirements.

CAD-bench evaluates complete systems rather than one modeling API or output format. A system receives a task prompt and produces either CAD code or a STEP artifact. The benchmark executes the submission and runs a grader written for that task. The 17 tasks cover basic solids, feature-rich parts, standards-like components, and functional assemblies. Table 2 states exactly what each grader checks.

The central claim is simple: build success is not enough. In the reported agent rows, for example, the right-angle gearbox builds in 96.9% of attempts but averages only 4.0% on the full score. The benchmark separates “the program ran” from “the artifact met the task.”

Contributions. This paper makes four contributions.

- We introduce CAD-bench, a 17-task executable benchmark for language-model CAD agents, spanning primitive solids, feature-rich parts, standards-like components, threaded mating pairs, and functional transmissions.
- We define a harness protocol that supports both one-shot CAD-code generation and agent-produced STEP submissions, allowing model-only and tool-using systems to be evaluated through a shared artifact interface.
- We develop a task-specific scoring stack that combines build success, dimensional and pose checks, reference-geometry gates, thread-profile analysis, and Blender-based rigid-body simulation.
- We report complete baseline sweeps for standalone models and agent harnesses, together with scorer-validation cases, aggregation-sensitivity checks, repeated-run diagnostics, and reproducibility artifacts.

2 Related Work

CAD-bench builds on CAD generation, engineering-document understanding, and executable agent benchmarks.

CAD generation and CAD code benchmarks. Seff et al. [2020] introduce SketchGraphs, a large collection of CAD sketches represented as geometric constraint graphs. Ganin et al. [2021] formulate CAD sketch generation as a language modeling problem over serialized design programs, while Wu et al. [2021] model CAD construction histories as operation sequences. Willis et al. [2021] provide human-authored Fusion 360 design sequences and an environment for programmatic CAD reconstruction. More recently, Khan et al. [2024] introduce Text2CAD for generating sequential CAD models from natural-language prompts, and Alrashedy et al. [2024] introduce CADPrompt, a benchmark of natural-language CAD prompts paired with expert CAD code, along with a visual

75 verification loop. CAD-Coder [Guan et al., 2025] and Text-to-CadQuery [Xie and Ju, 2025] move
76 directly toward text-to-parametric-code generation, using executable CadQuery and geometric rewards
77 such as Chamfer distance. CadEval [Epoch AI, 2026] evaluates text-to-CAD systems with rendering
78 success, Chamfer and Hausdorff distances, and volume or consistency checks. CADSmith [Barkley
79 et al., 2026] is especially close in spirit: it uses execution repair plus programmatic OpenCASCADE
80 measurements and visual judging to refine CadQuery code. CAD Arena [CAD Arena, 2026] provides
81 an emerging public comparison grid for 20 prompts, but its current reported metric is valid executable
82 STL production.

83 These works establish CAD as a structured generation problem and show that executable geometric
84 feedback matters. CAD-bench differs by scoring submitted artifacts against task-level mechanical
85 requirements, including mating interfaces, thread observables, exact pose, and rigid-body function,
86 rather than construction-sequence, rendering, or point-cloud similarity.

87 **Engineering-document understanding.** Doris et al. [2024] introduce DesignQA for evaluating
88 multimodal understanding of engineering documentation. That work is complementary to CAD-
89 bench: understanding drawings and manuals is part of the input side of engineering intelligence,
90 while CAD-bench evaluates output-side artifact generation in a runnable CAD environment.

91 **Executable benchmarks for agents.** Recent work has shown the value of realistic, environment-
92 grounded evaluation for agents. AgentBench [Liu et al., 2023], GAIA [Mialon et al., 2023], Visu-
93 alWebArena [Koh et al., 2024], and OSWorld [Xie et al., 2024] benchmark agents in interactive
94 environments rather than static text tasks. In software engineering, SWE-bench [Jimenez et al., 2023]
95 and MLE-bench [Chan et al., 2024] show that end-to-end execution environments uncover failure
96 modes that synthetic benchmarks miss. RE-Bench [Wijk et al., 2024] and PaperBench [Starace et al.,
97 2025] extend this perspective to open-ended research engineering. CAD-bench adopts the same
98 benchmark philosophy for CAD: models should be evaluated in the artifact-producing environment
99 itself.

100 **Benchmark methodology.** BetterBench [Reuel et al., 2024] emphasizes best practices such as
101 transparent reporting, replicability, and careful benchmark lifecycle management, while LiveBench
102 [White et al., 2024] highlights the importance of contamination-resistant evaluation. CAD-bench
103 follows these principles through executable scoring, explicit artifact traces, and a design that makes
104 future task refreshes and extensions straightforward.

105 3 Motivation: Why CAD Needs an Executable Benchmark

106 CAD evaluation requires stronger checks than syntactic validity or visual plausibility. A model
107 can generate valid Python, export a watertight mesh, or render an object with the correct silhouette
108 while still producing the wrong dimensions, pose, feature placement, mating geometry, or assembly
109 behavior. In mechanical design, these local errors are often decisive.

110 Surface-level geometric similarity is also incomplete. Many CAD tasks are defined by constrained
111 interfaces rather than by global shape alone. A screw head with an incorrect socket, a thread that
112 cannot mate, a gear train that collides, or a gearbox that fails to transmit motion may remain visually
113 plausible under coarse rendering or point-cloud metrics. Such artifacts may look close to a reference
114 but fail the engineering role specified by the prompt.

115 CAD agents also operate in executable tool environments, where they must synthesize parametric
116 programs, reason over APIs and exchange formats, and satisfy exact toolchain semantics. CAD-bench
117 is therefore designed around task-specific observables rather than visual complexity alone: visually
118 simple tasks may require exact conventions or mating constraints, while visually complex tasks may
119 be straightforward once decomposed into deterministic features.

Table 1: CAD-bench task taxonomy. Difficulty is authored for benchmark design; as shown later, model difficulty is not strictly monotone in the same order.

Difficulty	Count	Example tasks	Main capability stress
Easy	3	cube, box, ring spacer	pose, units, simple solids
Medium	4	L-bracket, extrusion segment, slotted plate, hex nut blank	feature placement, boolean modeling, conventions
Hard	6	flange, bolted ring, stepped block, large slotted plate, spur gear, M3 socket-head screw	denser geometry, standards-like structure, gears, threads
Functional	4	gearbox, custom threaded pair, compound gearbox, right-angle compound gearbox	assembly reasoning, interface compatibility, rigid-body functionality



Figure 2: Reference artifacts from the exact frozen task bundle used in the submission: a pose-sensitive cube, an M3 socket-head screw with thread and socket checks, and a two-gear transmission evaluated by rigid-body simulation. Each panel shows fixed front, top, and isometric views; the views illustrate the tasks but do not determine the score.

4 Benchmark Design and Artifact Construction

4.1 Scope and task taxonomy

CAD-bench contains 17 tasks divided into four difficulty levels: easy, medium, hard, and functional. The repository keeps the historical internal label `insane` for the functional bucket, but the paper uses “functional” because that is the capability being measured.

The tasks are a purpose-built coverage set, not a random or statistically representative sample of CAD work. We included a task only when (i) its prompt states observable geometric or mechanical requirements, (ii) those requirements can be checked deterministically in the released Linux toolchain, (iii) a reference and at least one controlled variation can exercise the grader, and (iv) it adds a failure mode not already covered by a simpler task. The sequence deliberately moves from pose and dimensions to feature placement, standards-like geometry, mating interfaces, and simulated motion. We exclude surface styling, free-form industrial design, large assemblies, finite-element analysis, and manufacturing-process planning from v0. Table 1 summarizes this coverage, Figure 2 shows three actual frozen references, and Table 2 lists every task and its checks.

The benchmark intentionally mixes single-part and multi-part tasks. The early tasks test whether a model can synthesize basic parametric geometry. Middle tasks add feature density and standards-like details. The final tasks require reasoning about interfaces and physical behavior rather than about isolated solids alone.

4.2 Task packaging

Each task is stored as a prompt, metadata, expected values, and evaluator configuration. In the repository, tasks are specified through files such as `prompt.txt`, `task.toml`, and optional reference assets used for reproducibility. The current public task format includes reference programs in `gold.py`, but those are benchmark-side reference implementations rather than the definition of the task itself. The expected metadata includes task ID, difficulty, evaluator, and task-specific expected values. Public task payloads can be loaded from a local mirror or from a Hugging Face repository.

This packaging supports two important use cases: exact reproduction of published harness runs and evaluation of custom harnesses against the same task set and scoring logic.

The tasks are synthetic but engineering-grounded: they specify concrete dimensions, poses, mating interfaces, thread conventions, and gearbox behavior that a CAD artifact must satisfy. This construc-

Table 2: Complete task-to-grader map. “Reference gate” is a coarse missing/extra-volume penalty; task checks remain responsible for the measurements named below. Functional rows use part-aware Blender simulation.

Tier	Task	Checks applied to the submitted artifact	Reference gate
Easy	cube	side length, volume, center, required negative- z pose	yes
Easy	box	three side lengths, volume, center, required positive- z pose	yes
Easy	ring spacer	inner/outer diameter, height, volume, center and pose	yes
Medium	L-bracket	bounding box, thickness, two-hole diameter/placement, volume	yes
Medium	extrusion	profile extent, length, center and outer bore diameters, pose	no
Medium	slotted plate	plate size, slot length/width, four holes, volume and pose	yes
Medium	hex nut blank	across-flats width, thickness, bore, volume and pose	yes
Hard	flange	base/boss/bore diameters, heights, four-hole circle and pose	yes
Hard	bolted ring	inner/outer diameter, height, six-hole circle and pose	yes
Hard	stepped block	overall extents, stepped volume, four holes and pose	yes
Hard	large slotted plate	plate/slot dimensions, eight holes, volume and pose	yes
Hard	spur gear	tooth count, module, root/outer diameter, bore, thickness and pose	no
Hard	M3 socket screw	pose, body/head dimensions, socket, thread pitch/handedness and insertion	no
Functional	two-gear transmission	axle fit/pose, separate bodies, collision contact, output direction and 2:1 speed ratio	no
Functional	custom threaded pair	two bodies, major/root diameters, two-start pitch/lead, clearance and mating observables	no
Functional	compound gearbox	three axle fits, two meshing stages, body separation and output speed ratio	no
Functional	right-angle gearbox	orthogonal axle/gear placement, envelope, contacts, reverse direction and output ratio	no

tion is intended to make the benchmark sound as an evaluation artifact: the prompts define real CAD outcomes, the metadata records the expected observables, and the evaluators test submitted artifacts rather than checking whether a model reproduced a particular text program.

4.3 Runtime and submission modes

The benchmark runtime supports two submission styles. In one-shot mode, current code-oriented harnesses return CAD code inside a structured XML block, with the built-in implementation using Build123D [build123d contributors, 2026]. In step-based agent mode, an interactive harness writes the final CAD artifact to a designated STEP path inside the evaluation environment. In both cases, the runtime executes the submission in a containerized workspace and converts the result into a common scoring pipeline.

This paper reports complete standalone model runs separately from complete agent runs. The agent interface remains part of the benchmark runtime, but agent rows are not mixed into the standalone model table because they use interactive tool feedback and a different wall-clock budget.

4.4 Scoring and verification

Each task returns normalized metrics including `build_success`, `task_score`, `overall_score`, and a benchmark-facing reward:

$$\text{reward} = 0.05 \cdot \text{build_success} + 0.95 \cdot \text{overall_score}. \quad (1)$$

The 5% build term records that the system produced an artifact the toolchain could open; 95% comes from satisfying the task. It prevents a parser failure and a mechanically wrong but inspectable artifact from being identical in diagnostic output, without allowing build success to dominate the result.

At the benchmark level, CAD-bench first averages task `overall_score` within each difficulty bucket. Missing benchmark tasks are inserted as explicit zero-score rows for result and paper aggregation. The final score is the category-weighted mean

$$S = \frac{1S_{\text{easy}} + 2S_{\text{medium}} + 3S_{\text{hard}} + 4S_{\text{functional}}}{1 + 2 + 3 + 4}. \quad (2)$$

This means a zero on a difficulty bucket remains part of the denominator rather than silently dropping that category from the leaderboard. The 1:2:3:4 weighting is a benchmark-design choice rather than a statistical estimate of task importance. It prevents the 13 static tasks from dominating the four functional tasks, while still reporting every bucket separately. We do not use any aggregation rule for statistically significant pairwise model-ranking claims; Section 7 reports a small aggregation-sensitivity check.

The task evaluators are designed to match engineering reality. The benchmark uses:

Table 3: Representative scorer-calibration cases. Full is the CAD-bench overall_score. Build is a build-success-only ablation. Proxy is the available generic geometry proxy where one exists; gearbox rows are dominated by functional simulation rather than a generic geometry proxy.

Task	Case	Full	Build	Proxy
Cube	independent valid cube	1.000	1.000	1.000
Cube	correct cube shifted off the required pose	0.000	1.000	0.000
Flange	independently constructed valid flange	1.000	1.000	1.000
Flange	bolt circle 1 mm too small	0.800	1.000	0.800
Flange	missing all four bolt holes	0.176	1.000	0.341
M3 screw	independently authored equivalent	1.000	1.000	1.000
M3 screw	threaded screw with missing socket	0.629	1.000	0.723
M3 screw	plain unthreaded shank	0.273	1.000	0.435
Gearbox	reference functional gear pair	0.962	1.000	–
Gearbox	rigid bridge between axle holes	0.000	1.000	–
Gearbox	reference gear pair shifted off axles	0.000	1.000	–
Threaded pair	reference mating threaded pair	0.974	1.000	–
Threaded pair	plausible unthreaded bolt and nut	0.406	1.000	–
Compound gearbox	rigid three-axle bridge	0.000	1.000	–
Right-angle gearbox	intersecting reference-like assembly	0.000	1.000	–

- mesh- and part-level dimensional checks,
- pose and placement metrics,
- reference-geometry gates that downweight large deviations from the authored solution,
- thread-specific profile and pitch metrics for the custom threaded-pair task, and
- Blender-based rigid-body simulation for gearbox tasks.

Reference-geometry gates are coarse missing/extra-volume checks used on selected static tasks. They do not compare modeling histories or rendered pixels. A gate multiplies the task score by a smooth penalty when the submitted shape differs greatly in occupied volume from the reference; the named task checks still score dimensions, pose, threads, and function. This catches artifacts that satisfy a few scalar measurements while omitting most of the object. It can also reject an unanticipated valid shape, so we audit it explicitly below.

This combination is central to the benchmark. A large part of CAD-bench’s value is that it measures the difference between “the program ran” and “the artifact solved the task.”

4.5 Scorer validation and ablations

Because CAD-bench uses task-specific evaluators, the release includes explicit scorer-calibration cases in `metadata/cad-bench-scorer-validation.json`. These cases cover reference submissions, independently authored valid submissions, minor defects, major defects, and nonbuild submissions across representative easy, hard-static, standards-like, threaded-pair, and functional tasks. Table 3 summarizes the main cases.

The validation cases serve two purposes. First, they test invariance to implementation details by giving full credit to independently authored valid artifacts that are not copies of the corresponding `gold.py` programs. Second, they test sensitivity to engineering-relevant defects. Several defective artifacts build successfully but receive low full scores, which demonstrates why build success alone is an inadequate CAD metric. Generic geometry proxies also over-credit local-but-incomplete artifacts, such as a screw without a socket or a flange without bolt holes. For gearbox tasks, the decisive signal is rigid-body behavior: a rigid bridge with plausible axle holes builds and may have the approximately correct bounding box, but receives zero because it cannot transmit the required motion.

Full-task scorer audit. We additionally constructed one independently expressed intended-valid artifact and one near-correct fault for each of the 17 tasks, including all four simulation tasks. The 17 references average 0.992, the intended-valid alternatives average 0.896, and the near-correct faults average 0.684. This broader audit found both false rejection and under-penalized defects. A valid hex nut rotated about its axis passes every named task check but scores zero because the reference-volume gate encodes an orientation that the prompt does not require. Conversely, a 0.5 mm extrusion-bore error scores 1.000, a narrowed gear tooth profile scores 0.943, and a 5% threaded-pair lead error scores 0.971. These are grader limitations, not model successes. The benchmark should therefore be

Table 4: Summary of complete CAD-bench baselines. Scores are category-weighted using Eq. (2). Full per-row results and run identifiers are reported in Appendix B.

Setting	Rows	Best overall	Mean easy	Mean hard	Mean functional
Standalone models	19	0.599	0.930	0.585	0.037
Agent harnesses	32	0.677	0.906	0.761	0.141

read with per-task diagnostics, and the hex-nut gate should be removed in the next task-set version rather than silently changed in the submitted version.

Worked M3 example. The M3 socket-head screw grader illustrates the scoring pattern in concrete terms. It checks that the head ends at $z = 0$ and the body points downward; measures the head, shank, and under-head length; probes the hexagonal socket; and estimates the thread’s pitch, handedness, and turns to seat. These measurements are combined into one score. An independently authored equivalent screw receives full credit, a threaded screw without a socket receives partial credit, and an unthreaded cylinder receives little credit even though all three build.

5 Experimental Protocol

5.1 Model and harness matrix

The benchmark is intended to support a mixed portfolio of closed and open providers. The current reviewer result payload contains 19 complete standalone model sweeps and 32 complete agent rows satisfying the validity criteria in Appendix B.3. Table 4 summarizes these rows in the main paper, while Appendix B records the exact run identifiers and full row-level results. The current standalone model set covers OpenAI, DeepSeek, Gemini, Vercel AI Gateway, and free OpenRouter rows.

The paper focuses on completed full-benchmark runs, meaning sweeps that cover all 17 tasks and reach model output plus benchmark scoring. This choice avoids mixing partially completed campaigns with full evaluations. Provider quota failures, API outages, malformed transport responses, harness setup failures, and upload failures are retained in the provenance artifacts for accounting but excluded from model-quality comparisons and from the reviewer result table because they do not measure the model’s CAD capability.

5.2 Reporting protocol

For each completed run, we report benchmark score, per-difficulty aggregates, wall-clock time, and estimated cost when the harness exposes enough usage and pricing data. When multiple complete runs exist for the same provider, model, access mode, and reasoning level, all complete rows in the reviewer result payload are reported, and the exact run identifiers are recorded in Appendix B. The current paper does not claim statistically significant pairwise rankings.

This choice reflects the current state of the experimental log rather than an intended limitation of the benchmark. CAD-bench itself stores the artifacts needed for repeated-run reporting, and larger future campaigns can replace unreplicated full-sweep rows with repeated-run aggregates.

6 Results: What CAD-bench Reveals

6.1 Overall performance

Table 4 summarizes the complete full-benchmark evaluations mirrored for anonymous review. Full leaderboards, run identifiers, wall-clock times, and cost estimates are reported in Appendix B; the main text focuses on aggregate conclusions because many configurations currently have only one complete sweep.

The 51 complete sweeps contain 867 task outcomes. Table 5 separates failures to produce geometry from artifacts that build but do not satisfy the task. The bins are descriptive score ranges, not new grading rules.

Table 5: Outcome audit over all 867 task attempts in the reported complete sweeps.

Observed outcome	Count	Fraction
Built, high score	506	58.4%
Build failure	126	14.5%
Built, zero or floor score	108	12.5%
Built, partial score	82	9.5%
Built, low score	45	5.2%

Three conclusions emerge from the completed rows.

First, CAD-bench is not saturated. The best current standalone model reaches 0.599 overall, and the best agent harness reaches 0.677. These scores leave substantial headroom despite high performance on many easy and static tasks.

Second, static geometry is substantially easier than functional CAD. Completed standalone runs average 0.585 on hard-static tasks but only 0.037 on functional tasks. Agent harnesses improve both categories, but the same pattern remains: they average 0.761 on hard-static tasks and 0.141 on functional tasks.

Third, the benchmark distinguishes model and harness settings. Agent harnesses generally outperform standalone one-shot generation, consistent with the hypothesis that execution feedback helps repair dimensional, geometric, and interface errors. However, agent success is uneven, and functional assemblies remain difficult even when submissions build successfully. The largest build-to-score gaps occur on the right-angle gearbox (0.667 build rate versus 0.025 mean score over all reported attempts), threaded pair (0.667 versus 0.049), compound gearbox (0.686 versus 0.112), and two-gear transmission (0.765 versus 0.223).

The difficulty claim is not based only on the names of the tiers. Under the same reported artifacts, replacing CAD-bench’s full score with weaker proxies would make the benchmark look easier. Reported agent submissions build the right-angle gearbox in 96.9% of attempts but average only 4.0% full score, and standalone submissions build the gearbox task in 63.2% of attempts but average only 6.7% full score. Thus the functional tasks are not merely harder by our label, they are clearly also harder for the language models.

6.2 Per-difficulty behavior

Across the 19 reported standalone model sweeps, the mean score by difficulty is 0.930 on easy, 0.753 on medium, 0.585 on hard, and 0.037 on functional tasks. Across the 32 reported agent sweeps, the corresponding means are 0.906, 0.849, 0.761, and 0.141. The drop on functional tasks is expected. The remaining drop from medium to hard indicates that standards-like details, dense feature placement, and gear/thread structure still matter after simple part synthesis is solved.

In CAD-bench, difficulty labels are benchmark-design labels, not claims about a strict monotone model ordering. Several hard tasks are still static, deterministic solids with many dimensions but little ambiguity, while some lower-level tasks are sensitive to conventions such as exact orientation, alignment, or hole placement. The current results therefore reinforce the point that CAD evaluation must be task-specific rather than purely complexity-labeled.

7 Discussion and Limitations

CAD-bench is mainly an evaluation artifact rather than a final model leaderboard. Its main claim is that executable CAD scoring exposes failures that build success, rendering validity, and coarse geometry proxies miss. The reported rows are useful diagnostic baselines, but many configurations have only one full sweep. We therefore avoid statistically significant pairwise model-ranking claims. Repeated-run statistics in Appendix B.2 support the static-versus-functional difficulty contrast for seven standalone configurations, and future leaderboard versions should promote rows only after repeated full sweeps under a declared aggregation rule.

Table 7 shows exactly how the authored 1:2:3:4 tier weights affect the leading rows. The top standalone row is unchanged under all three natural aggregations. The top agent row changes when

Table 6: Per-task mean overall_score and build success across the complete reported rows. Standalone means use 19 model sweeps; agent means use 32 agent sweeps. This table is computed from the 867 task rows in the full reviewer artifact, not from the difficulty aggregates alone.

Difficulty	Task	Standalone score	Standalone build	Agent score	Agent build
Easy	cube	0.895	0.947	0.938	0.969
Easy	box	0.947	1.000	0.844	0.844
Easy	ring spacer	0.947	1.000	0.938	0.938
Medium	L-bracket	0.632	0.842	0.906	0.969
Medium	extrusion	0.868	0.947	0.875	0.875
Medium	slotted plate	0.672	0.947	0.864	0.969
Medium	hex nut blank	0.842	0.895	0.750	0.906
Hard	flange	0.746	0.842	0.906	0.969
Hard	bolted ring	0.789	0.789	0.947	1.000
Hard	stepped block	0.812	0.947	1.000	1.000
Hard	large slotted plate	0.718	0.842	0.853	1.000
Hard	spur gear	0.237	0.526	0.345	0.844
Hard	M3 socket screw	0.207	0.526	0.512	0.875
Functional	gearbox	0.067	0.632	0.316	0.844
Functional	threaded pair	0.012	0.211	0.071	0.938
Functional	compound gearbox	0.069	0.632	0.138	0.719
Functional	right-angle gearbox	0.000	0.158	0.040	0.969

Table 7: Aggregation sensitivity. Scores shown are under the aggregation named in each row.

Aggregation	Top standalone	Score	Top agent	Score
Authorized tier weights (1:2:3:4)	Gemini 3.1 Pro offline	0.599	Codex GPT-5.4 mini web medium	0.677
Equal weight per tier	Gemini 3.1 Pro offline	0.727	Codex GPT-5.4 mini web medium	0.781
Equal weight per task	Gemini 3.1 Pro offline	0.709	Pi GPT-5.4 offline high	0.772

every task, rather than every tier, receives equal weight. Leave-one-task-out analysis similarly gives the weighted top standalone row 16/17 wins, but the weighted top agent row only 12/17. A task bootstrap gives the top standalone row 0.645 probability of ranking first and the top agent row 0.255. These are task-composition checks, not repeated-run confidence intervals. The qualitative static-versus-functional gap is stable, but the exact agent ordering is not.

The benchmark is also small. Seventeen tasks are enough for a high-information initial release that covers primitives, feature placement, standards-like details, threaded interfaces, and simulated mechanisms, but they are not enough to claim broad CAD-agent competence. The public release is therefore intended to be versioned. Appendix A describes the planned lifecycle: immutable public task bundles, private refresh tasks with canary hashes, contamination monitoring, and repeated-run promotion for leaderboard rows.

Task-specific scorers encode engineering assumptions. This is necessary for CAD: a threaded pair, socket-head screw, or gearbox cannot be judged by build success alone. It also creates a risk that a valid alternative solution could miss an authored observable. The scorer-validation cases in Table 3 partially address this by giving full credit to independently authored valid artifacts and low credit to controlled defects, but they are not exhaustive. The intended use of CAD-bench is therefore comparative and diagnostic: inspect per-task traces, scorer components, and validation cases, not only the aggregate score.

8 Conclusion

CAD-bench is an executable benchmark for CAD generation that emphasizes artifact correctness over surface plausibility. The current v0 release shows that models and agent harnesses can often produce runnable CAD, but still fail on standards-like details, mating interfaces, and functional assemblies. The benchmark’s main contribution is the scoring interface and reproducible artifact protocol: it separates easy geometry from mechanically meaningful CAD behavior and gives future systems a concrete target for improving assemblies and interfaces that actually work.

References

- Kamel Alrashedy, Pradyumna Tambwekar, Zulfiqar Zaidi, Megan Langwasser, Wei Xu, and Matthew Gombolay. Generating CAD Code with Vision-Language Models for 3D Designs. arXiv:2410.05340, 2024. URL: <https://arxiv.org/abs/2410.05340>.
- Jesse Barkley, Rumi Loghmani, and Amir Barati Farimani. CADSmith: Multi-Agent CAD Generation with Programmatic Geometric Validation. arXiv:2603.26512, 2026. URL: <https://arxiv.org/abs/2603.26512>.
- build123d contributors. build123d documentation. 2026. URL: <https://build123d.readthedocs.io/en/stable/index.html>.
- CAD Arena. CAD Arena: Open Benchmark for AI-Generated Parametric CAD. 2026. URL: <https://cadarena.dev/>.
- Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, James Aung, Dane Sherburn, Evan Mays, Giulio Starace, Kevin Liu, Leon Maksin, Tejal Patwardhan, Lilian Weng, and Aleksander Mądry. MLE-bench: Evaluating Machine Learning Agents on Machine Learning Engineering. arXiv:2410.07095, 2024. URL: <https://arxiv.org/abs/2410.07095>.
- Anna C. Doris, Daniele Grandi, Ryan Tomich, Md Ferdous Alam, Mohammadmehdi Ataei, Hyunmin Cheong, and Faez Ahmed. DesignQA: A Multimodal Benchmark for Evaluating Large Language Models’ Understanding of Engineering Documentation. arXiv:2404.07917, 2024. URL: <https://arxiv.org/abs/2404.07917>.
- Epoch AI. CadEval. 2026. URL: <https://epoch.ai/benchmarks/cad-eval>.
- Yaroslav Ganin, Sergey Bartunov, Yujia Li, Ethan Keller, and Stefano Saliceti. Computer-Aided Design as Language. arXiv:2105.02769, 2021. URL: <https://arxiv.org/abs/2105.02769>.
- Yandong Guan, Xilin Wang, Xingxi Ming, Jing Zhang, Dong Xu, and Qian Yu. CAD-Coder: Text-to-CAD Generation with Chain-of-Thought and Geometric Reward. arXiv:2505.19713, 2025. URL: <https://arxiv.org/abs/2505.19713>.
- Mohammad Sadil Khan, Sankalp Sinha, Talha Uddin Sheikh, Didier Stricker, Sk Aziz Ali, and Muhammad Zeshan Afzal. Text2CAD: Generating Sequential CAD Models from Beginner-to-Expert Level Text Prompts. In *Advances in Neural Information Processing Systems*, volume 37, 2024. URL: <https://openreview.net/forum?id=5k9XeHIK3L>.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? arXiv:2310.06770, 2023. URL: <https://arxiv.org/abs/2310.06770>.
- Jing Yu Koh, Robert Lo, Lawrence Jang, Vikram Duvvur, Ming Chong Lim, Po-Yu Huang, Graham Neubig, Shuyan Zhou, Ruslan Salakhutdinov, and Daniel Fried. VisualWebArena: Evaluating Multimodal Agents on Realistic Visual Web Tasks. arXiv:2401.13649, 2024. URL: <https://arxiv.org/abs/2401.13649>.
- Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. AgentBench: Evaluating LLMs as Agents. arXiv:2308.03688, 2023. URL: <https://arxiv.org/abs/2308.03688>.
- Gregoire Mialon, Clementine Fourier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. GAIA: a benchmark for General AI Assistants. arXiv:2311.12983, 2023. URL: <https://arxiv.org/abs/2311.12983>.
- Anka Reuel, Amelia Hardy, Chandler Smith, Max Lamparth, Malcolm Hardy, and Mykel J. Kochenderfer. BetterBench: Assessing AI Benchmarks, Uncovering Issues, and Establishing Best Practices. arXiv:2411.12990, 2024. URL: <https://arxiv.org/abs/2411.12990>.

- 366 Ari Seff, Yaniv Ovadia, Wenda Zhou, and Ryan P. Adams. SketchGraphs: A Large-Scale Dataset
367 for Modeling Relational Geometry in Computer-Aided Design. arXiv:2007.08506, 2020. URL:
368 <https://arxiv.org/abs/2007.08506>.
- 369 Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias,
370 Evan Mays, Benjamin Kinsella, Wyatt Thompson, Johannes Heidecke, Amelia Glaese, and Tejal
371 Patwardhan. PaperBench: Evaluating AI’s Ability to Replicate AI Research. arXiv:2504.01848,
372 2025. URL: <https://arxiv.org/abs/2504.01848>.
- 373 Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Ben Feuer, Siddhartha Jain, Ravid Shwartz-
374 Ziv, Neel Jain, Khalid Saifullah, Siddhartha Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein,
375 Willie Neiswanger, and Micah Goldblum. LiveBench: A Challenging, Contamination-Free LLM
376 Benchmark. arXiv:2406.19314, 2024. URL: <https://arxiv.org/abs/2406.19314>.
- 377 Haoyang Xie and Feng Ju. Text-to-CadQuery: A New Paradigm for CAD Generation with Scalable
378 Large Model Capabilities. arXiv:2505.06507, 2025. URL: <https://arxiv.org/abs/2505.06507>.
- 380 Karl D. D. Willis, Yewen Pu, Jieliang Luo, Hang Chu, Tao Du, Joseph G. Lambourne, Armando
381 Solar-Lezama, and Wojciech Matusik. Fusion 360 Gallery: A Dataset and Environment for
382 Programmatic CAD Construction from Human Design Sequences. *ACM Transactions on Graphics*,
383 40(4), 2021. URL: <https://arxiv.org/abs/2010.02392>.
- 384 Hjalmar Wijk, Tao Lin, Joel Becker, Sami Jawhar, Neev Parikh, Thomas Broadley, Lawrence Chan,
385 Michael Chen, Josh Clymer, Jai Dhyani, Elena Elicheva, Katharyn Garcia, Brian Goodrich,
386 Nikola Jurkovic, Megan Kinniment, Aron Lajko, Seraphina Nix, Lucas Sato, William Saunders,
387 Maksym Taran, Ben West, and Elizabeth Barnes. RE-Bench: Evaluating frontier AI R&D
388 capabilities of language model agents against human experts. arXiv:2411.15114, 2024. URL:
389 <https://arxiv.org/abs/2411.15114>.
- 390 Rundi Wu, Chang Xiao, and Changxi Zheng. DeepCAD: A Deep Generative Network for Computer-
391 Aided Design Models. In *Proceedings of the IEEE/CVF International Conference on Computer*
392 *Vision*, 2021. URL: <https://arxiv.org/abs/2105.09492>.
- 393 Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing
394 Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, Yitao Liu, Yiheng Xu, Shuyan Zhou, Silvio
395 Savarese, Caiming Xiong, Victor Zhong, and Tao Yu. OSWorld: Benchmarking Multimodal
396 Agents for Open-Ended Tasks in Real Computer Environments. arXiv:2404.07972, 2024. URL:
397 <https://arxiv.org/abs/2404.07972>.

398 A Benchmark lifecycle, release policy, and broader impacts

399 The artifact is designed to support a versioned benchmark lifecycle rather than a single frozen
400 leaderboard. The submitted release should be interpreted as CAD-bench v0: a public, reproducible
401 task set with exact run artifacts. Future leaderboard versions should use four controls.

- 402 • **Frozen public splits.** Published task bundles are immutable and identified by manifest
403 hashes and dataset revisions.
- 404 • **Private refresh tasks.** New functional and standards-like tasks can be staged privately, with
405 salted bundle-hash canaries exported before evaluation to establish that the tasks existed
406 before result collection.
- 407 • **Versioned leaderboards.** Rows should be reported against an explicit task-set version; old
408 rows should not be silently mixed with refreshed tasks.
- 409 • **Repeated-run promotion.** A configuration should move from diagnostic reporting to a
410 primary leaderboard row only after independent full sweeps are available with a declared
411 aggregation rule.

412 This policy is already reflected in the release files: task manifests record per-file hashes, the release
413 pins the task dataset revision, and the source artifact contains a canary-export command for private

or unreleased tasks. The current paper still reports v0 diagnostic rows because the full repeated-run campaign is not complete; this is why the claims are framed around benchmark behavior and failure modes rather than statistically significant pairwise model rankings.

CAD agents could provide several benefits: lowering the barrier to prototyping, supporting accessibility-focused device customization, improving engineering education, and reducing repetitive modeling work in low-volume manufacturing workflows.

The same capabilities also create risks. More capable CAD generation could accelerate the design of unsafe hardware, restricted components, or deceptive engineering artifacts. It may also encourage overtrust in automatically generated mechanical designs that have not received appropriate engineering review. Transparent, function-aware benchmarking is one mitigation: it makes it harder to confuse runnable CAD code with mechanically correct design.

B Reproducibility and Reporting Details

B.1 Raw complete result rows

Tables 8–10 record the exact complete rows summarized in Table 4. Scores are reported as percentages for readability. The run column is the run identifier corresponding to the stored report artifact. Partial runs and infrastructure-failed rows are excluded from model-quality comparisons.

Table 8: Raw complete standalone model rows.

Rank	Harness	Overall	Easy	Med.	Hard	Func.	Min.	Cost	Run
1	gemini/gemini-3.1-pro-preview-offline-none	59.9	100.0	96.8	70.8	23.2	25.4	5.8625	gemini__gemini-3.1-pro-preview-offline-none_20260428T194054Z_2313253
2	openai/gpt-5.4-pro-web-med	53.9	100.0	96.8	78.5	2.5	325.7	37.0405	openai__gpt-5.4-pro-web-med_20260423T122643Z_1626647
3	vercel/alibaba/qwen-3.6-max-preview-offline-none	52.8	100.0	96.8	76.2	1.5	42.8	3.2293	vercel__alibaba__qwen-3.6-max-preview-offline-none_20260430T003706Z_2473332
4	gemini/gemini-3-flash-preview-offline-none	52.6	100.0	96.8	77.4	0.0	21.1	1.3596	gemini__gemini-3-flash-preview-offline-none_20260428T121645Z_2280737
5	openrouter/tencent/hy3-preview-free-offline-none	51.1	100.0	96.8	60.8	8.7	51.4	—	openrouter__tencent__hy3-preview-free-offline-none_20260428T035110Z_2240039
6	openai/gpt-5.3-codex-spark-offline-med	50.4	100.0	96.8	70.0	0.0	6.1	4.9685	openai__gpt-5.3-codex-spark-offline-med_20260427T023421Z_2115684
7	vercel/alibaba/qwen3.6-27b-offline-none	50.0	100.0	96.8	68.9	0.0	58.6	1.5128	vercel__alibaba__qwen3.6-27b-offline-none_20260430T003707Z_2473345
8	openai/gpt-5.3-codex-spark-offline-xhigh	48.9	100.0	96.8	65.0	0.0	6.0	5.4877	openai__gpt-5.3-codex-spark-offline-xhigh_20260427T024838Z_2122296
9	openai/gpt-5.5-offline-med	48.7	100.0	71.8	65.0	12.0	9.8	8.6697	openai__codex__gpt-5.5-offline-med_20260426T055852Z_2010916
10	openai/gpt-5.5-offline-low	47.5	100.0	71.8	76.3	0.6	10.1	9.7384	openai__codex__gpt-5.5-offline-low_20260426T054800Z_2000292
11	deepseek/deepseek-v4-pro-offline-none	47.5	100.0	59.4	68.0	12.9	55.7	3.6274	deepseek__deepseek-v4-pro-offline-none_20260426T042645Z_1958729
12	openai/gpt-5.5-offline-high	45.6	66.7	71.8	82.1	0.0	8.9	7.8075	openai__codex__gpt-5.5-offline-high_20260426T061252Z_2019625
13	openai/gpt-5.5-offline-xhigh	45.6	100.0	71.8	59.3	8.7	9.9	9.2523	openai__codex__gpt-5.5-offline-xhigh_20260426T062226Z_2023449
14	openai/gpt-5.3-codex-spark-offline-high	40.4	100.0	96.8	36.9	0.0	6.1	5.2469	openai__gpt-5.3-codex-spark-offline-high_20260427T024046Z_2118961
15	deepseek/deepseek-v4-flash-offline-none	39.5	100.0	64.3	55.5	0.0	18.1	0.2778	deepseek__deepseek-v4-flash-offline-none_20260426T052907Z_1985868
16	gemini/gemini-3.1-flash-lite-preview-offline-none	34.2	100.0	75.0	30.8	0.0	19.2	0.1993	gemini__gemini-3.1-flash-lite-preview-offline-none_20260428T050032Z_2249465
17	openai/gpt-5.3-codex-spark-offline-low	33.9	100.0	71.8	31.7	0.0	5.8	4.2688	openai__gpt-5.3-codex-spark-offline-low_20260427T022715Z_2111735
18	openrouter/inclusionai/ling-2.6-1t-free-offline-none	13.9	33.3	2.7	33.3	0.0	12.8	—	openrouter__inclusionai__ling-2.6-1t-free-offline-none_20260428T033043Z_2234914
19	gemini/gemma-3-27b-it-offline-nodocs-none	8.1	66.7	0.0	4.6	0.0	3.7	—	gemini__gemma-3-27b-it-offline-nodocs-none_20260430T022138Z_2493048

B.2 Repeated-run statistical check

The current paper reports complete full 17-task sweeps for every row, but many configurations have only one independent sweep. The paper therefore does not use those rows for pairwise model-ranking significance claims. To test the main qualitative difficulty claim under repeated sampling, the source artifact includes 47 complete local reports, with repeated full-sweep statistics for seven standalone configurations, summarized in metadata/cad-bench-repeat-stats.json. Additional attempted repeats across the reported standalone and agent matrix were excluded when

Table 9: Raw complete agent rows.

Rank	Agent Harness	Overall	Easy	Med.	Hard	Func.	Min.	Cost	Run
1	Codex codex/gpt-5.4-mini-web-med	67.7	100.0	96.8	78.7	36.8	234.0	2.4446	codex_gpt-5.4-mini-web-med_20260420T00623Z_676664
2	Pi pi/gpt-5.4-offline-high	67.1	100.0	96.8	82.9	32.1	618.8	4.6416	pi_gpt-5.4-offline-high_20260422T092300Z_1340452
3	Codex codex/gpt-5.4-offline-high	65.3	66.7	71.8	83.6	47.9	245.2	8.2311	codex_gpt-5.4-offline-high_20260419T190624Z_603644
4	Pi pi/gpt-5.4-offline-xhigh	64.7	100.0	96.8	88.0	22.3	418.3	4.8298	pi_gpt-5.4-offline-xhigh_20260421T075050Z_1041869
5	Pi pi/gpt-5.4-web-high	64.6	100.0	96.8	82.9	25.9	281.6	4.5650	pi_gpt-5.4-web-high_20260422T092300Z_1340450
6	Pi pi/gpt-5.4-mini-web-xhigh	61.0	100.0	100.0	84.5	14.2	240.2	2.4442	pi_gpt-5.4-mini-web-xhigh_20260422T092301Z_1340454
7	Pi pi/gpt-5.4-mini-offline-high	60.3	100.0	75.0	80.3	28.0	169.0	3.5843	pi_gpt-5.4-mini-offline-high_20260422T151655Z_1441926
8	Pi pi/gpt-5.4-offline-med	59.9	100.0	96.8	76.2	19.3	304.2	1.9074	pi_gpt-5.4-offline-med_20260421T075052Z_1041872
9	Codex codex/gpt-5.4-web-low	59.1	100.0	96.8	79.4	14.8	279.8	6.3453	codex_gpt-5.4-web-low_20260419T190502Z_603630
10	Codex codex/gpt-5.4-web-med	59.1	100.0	71.8	78.0	28.2	260.7	5.7144	codex_gpt-5.4-web-med_20260419T190540Z_603629
11	Pi pi/gpt-5.4-mini-offline-xhigh	59.0	100.0	96.8	83.5	11.4	282.4	3.4776	pi_gpt-5.4-mini-offline-xhigh_20260422T121627Z_1402824
12	Codex codex/gpt-5.4-mini-offline-med	57.8	100.0	96.8	52.5	31.7	157.2	3.3316	codex_gpt-5.4-mini-offline-med_20260420T004107Z_683812
13	Codex codex/gpt-5.4-web-xhigh	56.9	100.0	96.8	89.7	1.7	396.2	12.7137	codex_gpt-5.4-web-xhigh_20260421T075051Z_1041874
14	Codex codex/gpt-5.4-mini-offline-xhigh	56.9	100.0	96.8	77.8	10.5	120.9	4.0830	codex_gpt-5.4-mini-offline-xhigh_20260422T05507Z_1497871
15	Pi pi/gpt-5.4-web-xhigh	56.0	100.0	71.8	82.9	16.9	419.3	4.7572	pi_gpt-5.4-web-xhigh_20260421T075052Z_1041870
16	Pi pi/gpt-5.4-mini-web-med	54.5	100.0	96.8	81.8	1.5	191.2	1.6381	pi_gpt-5.4-mini-web-med_20260422T092300Z_1340457

Table 10: Raw complete agent rows, continued.

Rank	Agent Harness	Overall	Easy	Med.	Hard	Func.	Min.	Cost	Run
17	Pi pi/gpt-5.4-mini-web-high	54.0	100.0	96.8	75.7	4.7	130.5	2.5139	pi_gpt-5.4-mini-web-high_20260422T183318Z_1468372
18	Codex codex/gpt-5.4-mini-web-high	53.8	100.0	96.8	78.3	2.3	389.9	3.6274	codex_gpt-5.4-mini-web-high_20260420T235625Z_845623
19	Codex codex/gpt-5.4-offline-xhigh	53.5	33.3	71.8	76.2	32.3	289.6	11.4765	codex_gpt-5.4-offline-xhigh_20260419T190535Z_603631
20	Pi pi/gpt-5.4-offline-low	53.3	100.0	75.0	84.4	7.6	171.2	1.5744	pi_gpt-5.4-offline-low_20260419T061026Z_383966
21	Pi pi/gpt-5.4-web-med	53.0	100.0	71.8	80.1	11.6	366.2	2.1871	pi_gpt-5.4-web-med_20260421T075053Z_1041868
22	Pi pi/gpt-5.4-web-low	52.6	100.0	96.8	76.2	0.9	65.1	1.1919	pi_gpt-5.4-web-low_20260422T205357Z_1497569
23	Codex codex/gpt-5.4-web-high	52.1	66.7	96.8	80.1	5.0	300.5	7.9083	codex_gpt-5.4-web-high_20260419T190546Z_603612
24	Codex codex/gpt-5.4-mini-offline-high	51.3	100.0	96.8	71.1	1.5	169.1	4.4820	codex_gpt-5.4-mini-offline-high_20260420T004018Z_683339
25	Codex codex/gpt-5.4-mini-web-xhigh	51.1	100.0	71.8	86.9	1.6	416.0	5.9764	codex_gpt-5.4-mini-web-xhigh_20260421T075051Z_1041875
26	Pi pi/gpt-5.4-mini-offline-med	49.7	100.0	96.8	64.5	2.4	241.8	2.4102	pi_gpt-5.4-mini-offline-med_20260422T094713Z_1351194
27	Pi pi/gpt-5.4-mini-web-low	48.1	100.0	75.0	77.1	0.0	189.4	0.5317	pi_gpt-5.4-mini-web-low_20260419T064631Z_413438
28	Codex codex/gpt-5.4-mini-offline-low	48.1	100.0	100.0	57.5	2.2	111.3	1.6688	codex_gpt-5.4-mini-offline-low_20260420T014421Z_709021
29	Codex codex/gpt-5.4-offline-low	46.8	33.3	46.8	88.4	18.9	181.1	4.4749	codex_gpt-5.4-offline-low_20260419T190806Z_603646
30	Pi pi/gpt-5.4-mini-offline-low	37.9	100.0	71.8	38.5	4.9	132.1	0.5499	pi_gpt-5.4-mini-offline-low_20260422T100629Z_1359611
31	Codex codex/gpt-5.4-offline-med	36.2	33.3	46.8	61.6	12.6	260.4	8.0969	codex_gpt-5.4-offline-med_20260419T190604Z_603645
32	Codex codex/gpt-5.4-mini-web-low	32.9	66.7	50.0	54.1	0.0	246.5	1.8227	codex_gpt-5.4-mini-web-low_20260420T000933Z_677450

437 provider quota, authentication, unavailable-model, or timeout failures prevented a complete 17-task
 438 report.

439 Across six openai/gpt-5.3-codex-spark-offline-med sweeps, the benchmark score is 0.471
 440 with bootstrap 95% CI [0.453, 0.490]. The easy bucket is stable at 1.000, medium is 0.973 with CI
 441 [0.968, 0.984], hard is 0.572 with CI [0.507, 0.643], and functional is 0.011 with CI [0.000, 0.033].
 442 A paired exact two-sided sign-flip test over full-sweep difficulty aggregates gives $p = 0.03125$ for
 443 hard score minus functional score, with mean difference 0.561 and bootstrap 95% CI [0.486, 0.641].

Across the four six-sweep `openai/gpt-5.5-offline-*` repeated checks, benchmark means are 0.477 for low with CI [0.455, 0.492], 0.444 for medium with CI [0.391, 0.503], 0.505 for high with CI [0.454, 0.563], and 0.573 for xhigh with CI [0.528, 0.612]. Their hard-minus-functional mean differences are 0.701, 0.483, 0.695, and 0.694, respectively, and each exact sign-flip test gives $p = 0.03125$. The xhigh row has the strongest repeated GPT-5.5 mean here, but these runs still support only a within-configuration static-versus-functional difficulty contrast, not a statistically significant model-ranking claim.

Across six `openrouter/tencent/hy3-preview:free-offline-none` sweeps, the benchmark score is 0.518 with CI [0.471, 0.557]; easy is 1.000, medium is 0.823 with CI [0.729, 0.917], hard is 0.664 with CI [0.585, 0.737], and functional is 0.135 with CI [0.095, 0.176]. The hard-minus-functional sign-flip test gives $p = 0.03125$, mean difference 0.529, and CI [0.472, 0.594]. Across five `openrouter/inclusionai/ling-2.6-1t:free-offline-none` sweeps, the benchmark score is 0.154 with CI [0.080, 0.236]; hard-minus-functional remains positive at 0.227 with CI [0.084, 0.407], but the exact sign-flip test is $p = 0.125$ at $n = 5$. These repeated-run checks support the paper’s main diagnostic claim that functional CAD remains much harder than static CAD for multiple standalone harnesses, while leaving the model-level ranking tables as diagnostic baselines rather than statistically significant pairwise comparisons.

The source artifact also retains the earlier `metadata/cad-bench-repeat-check.json` two-run sanity record for traceability. Future leaderboard versions should report repeated sweeps in the following format:

- number of independent full sweeps per configuration,
- central estimate (mean or median),
- uncertainty interval (e.g., standard deviation or bootstrap 95% CI),
- policy for selecting primary leaderboard rows when multiple repetitions exist.

The current analysis sections avoid pairwise significance claims and use the complete sweeps as baseline diagnostics.

B.3 Run validity criteria

A run is eligible for a complete-results table if it satisfies all of the following conditions:

- it uses a standalone model harness or an agent harness,
- it attempts the full 17-task benchmark under one fixed provider, model, access mode, and reasoning setting,
- each task reaches model-output capture, STEP export when applicable, and benchmark scoring, and
- failures, if any, are model-output or scored-artifact failures rather than provider quota, API outage, harness setup, or artifact-upload failures.

Provider and infrastructure failures are retained in the experiment logs for auditability, but they are not assigned model-quality scores because they do not evaluate the submitted model on CAD generation.

B.4 Compute and environment reporting

The benchmark runtime uses Dockerized environments, Python 3.11, NumPy, trimesh, and Blender for the functional simulation tasks, together with Build123D in the current code-oriented reference and harness pipeline. The reported local runs used the following environment:

- host hardware: Intel Core i5-7300U CPU, 15 GiB RAM, CPU-only scoring except for Blender’s local simulation/rendering stack,
- software versions: Docker 29.4.0, Blender 5.1.1, uv 0.11.7, project Python 3.11 environment from `uv.lock`,
- standalone model wall-clock range in the reported rows: 3.7–325.7 minutes,
- agent wall-clock range in the reported rows: 65.1–618.8 minutes,

- failed-run accounting: provider quota failures are kept in logs but excluded from model-quality comparisons.

B.5 Code, data, and asset availability

The benchmark runtime, scoring code, and harness interface are available in the anonymous source artifact accompanying this paper. Public task payloads are hosted in an anonymous reviewer data mirror, and the reviewer mirror includes a compact result artifact with every complete run reported in this paper.

For reviewer inspection, the artifact has a no-API smoke path: `uv run pytest tests/test_anonymous_artifacts.py tests/test_scorer_validation_artifact.py` checks the anonymous packaging and scorer-calibration artifact, and `uv run python scripts/export_scorer_validation.py` regenerates the calibration JSON. Full model sweeps require provider credentials, but scorer validation and artifact consistency do not.

The anonymous submission release pointers are:

- anonymized source archive: included in the supplemental material and mirrored with the anonymous dataset artifact,
- anonymized public task dataset mirror: included in the supplemental material,
- task dataset revision used for this paper: `20ad0b586b70ac2c5e6fe2b501386cf26a887137`,
- reviewer result artifact: `results/cad-bench-reported-results.json` in the anonymous dataset mirror,
- scorer validation artifact: `results/cad-bench-scorer-validation.json` in the anonymous dataset mirror,
- run artifact bundles: `provenance/<run-id>/report.json` entries in the anonymous reviewer artifact,
- provenance URLs: redacted from submitted metadata when they would reveal non-anonymous infrastructure.

The full reviewer artifact includes Croissant metadata with Responsible AI fields describing intended use, limitations, task authoring, and safety-relevant non-use cases. This matches the 2026 Evaluations & Datasets track expectation that datasets document how they support evaluative claims rather than serving only as raw files.

B.6 Licenses and upstream assets

Table 11 documents the main upstream software and data assets used by the benchmark runtime and reviewer dataset.

Table 11: Primary upstream software assets used by CAD-bench.

Asset	License / terms	Role in benchmark
Build123D	Apache-2.0	current reference implementation and code-harness dependency, not the benchmark definition itself
Blender	GNU GPL	rigid-body simulation and rendering for functional gearbox evaluation
NumPy	BSD-style license	numeric processing in evaluators and utilities
trimesh	MIT	mesh conversion and geometric analysis
Repository code	MIT	benchmark runtime, harnesses, and scoring stack
Public task dataset	MIT	hosted task prompts, metadata, reference programs, and authored benchmark assets
Authored CAD fixtures	MIT	synthetic task-side fixtures, including the M3 socket-head screw equivalent fixture
Closed model APIs	provider terms	external inference services used for benchmarked models

The benchmark tasks and authored fixtures are synthetic benchmark assets rather than redistributed proprietary CAD files.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [Yes]

Justification: The abstract and Sections 1 and 4 state the paper's scope as introducing an executable CAD benchmark, documenting its scoring stack, and reporting the currently completed baseline results with explicit limitations around unreplicated full-sweep comparisons.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The main results section and Appendix A discuss benchmark size, the current emphasis of the reported baselines on Build123D-based harnesses rather than the benchmark definition itself, incomplete all-provider campaign coverage, current dependence on a small number of completed full sweeps, benchmark exposure risk, and the planned refresh/holdout protocol.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is a benchmark and empirical evaluation paper and does not include theoretical results or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Sections 4 and 5 describe the evaluation protocol and reported metrics, and Appendix A records completed run identifiers, exact dataset revision, code/data availability, and compute details.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The paper describes the benchmark runtime, scoring code, and task packaging as releasable assets, and the appendix describes the anonymous source artifact, reviewer data mirror, pinned task revision, and result artifacts.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Sections 4 and 5 specify the relevant settings for this evaluation paper: provider, model, access mode, reasoning level, task coverage, benchmark score definition, per-difficulty reporting, and build-success aggregation.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The paper reports bootstrap confidence intervals and exact paired sign-flip tests for repeated full-sweep checks across seven standalone configurations where repeat counts support them. It still explicitly avoids pairwise model-ranking significance claims for configurations without enough repeated independent sweeps, and Appendix A states how repeated-run aggregates should be reported in future leaderboard versions.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix A reports host hardware, key software versions, full-sweep wall-clock ranges for standalone and agent rows, and how provider quota failures are accounted for.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work is a benchmark and evaluation study, does not involve deception or human subjects, and includes explicit discussion of limitations, failure modes, and broader impacts in Sections 6 and 7.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Limitations section and Appendix A discuss positive impacts for engineering productivity and accessibility, and negative impacts including unsafe hardware design, deceptive artifacts, and overtrust in generated designs.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases a benchmark runtime and evaluation assets, not a new high-risk generative model or scraped dataset that would require controlled release safeguards of the type described here.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The bibliography credits the main upstream research assets, and Appendix A includes a dedicated license table for the primary software dependencies, project code, public task dataset, authored fixtures, and provider terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The benchmark task taxonomy, runtime, scoring protocol, run accounting, and release structure are documented in Sections 4–6 and Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The paper does not involve crowdsourcing or research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

846 Answer: [N/A]
847 Justification: The paper does not involve human subjects research.
848 Guidelines:
849 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
850 with human subjects.
851 • Depending on the country in which research is conducted, IRB approval (or equivalent)
852 may be required for any human subjects research. If you obtained IRB approval, you
853 should clearly state this in the paper.
854 • We recognize that the procedures for this may vary significantly between institutions
855 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
856 guidelines for their institution.
857 • For initial submissions, do not include any information that would break anonymity (if
858 applicable), such as the institution conducting the review.

859 **16. Declaration of LLM usage**
860 Question: Does the paper describe the usage of LLMs if it is an important, original, or
861 non-standard component of the core methods in this research? Note that if the LLM is used
862 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
863 scientific rigor, or originality of the research, declaration is not required.

864 Answer: [N/A]
865 Justification: LLMs are not involved as any important, original, or non-standard components.
866 Guidelines:
867 • The answer [N/A] means that the core method development in this research does not
868 involve LLMs as any important, original, or non-standard components.
869 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
870 be described.